

## Airline Review Classification

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import spacy
from sklearn.feature_extraction.text import TfidfVectorizer
import tensorflow as tf
from tensorflow import keras
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import classification_report
import pickle
```

```
df=pd.read_csv(r"D:\MachineLearning\DataScienceCourse\
Airline_review.csv")
df.head(3)
```

	Unnamed: 0	Airline	Name	Overall_Rating	
Review_Title \					
0	0	AB Aviation	9	"pretty decent airline"	
1	1	AB Aviation	1	"Not a good airline"	
2	2	AB Aviation	1	"flight was fortunately short"	

	Review Date	Verified	\
0	11th November 2019	True	
1	25th June 2019	True	
2	25th June 2019	True	

	Review	Aircraft	\
0	Moroni to Moheli. Turned out to be a pretty ...	NaN	
1	Moroni to Anjouan. It is a very small airline...	E120	
2	Anjouan to Dzaoudzi. A very small airline an...	Embraer E120	

	Type Of Traveller	Seat Type	Route	Date Flown	\
0	Solo Leisure	Economy Class	Moroni to Moheli	Nov-19	
1	Solo Leisure	Economy Class	Moroni to Anjouan	Jun-19	
2	Solo Leisure	Economy Class	Anjouan to Dzaoudzi	Jun-19	

	Seat Comfort	Cabin Staff Service	Food & Beverages	Ground Service	\
--	--------------	---------------------	------------------	----------------	---

0	4.0	5.0	4.0	4.0
1	2.0	2.0	1.0	1.0
2	2.0	1.0	1.0	1.0

	Inflight Entertainment	Wifi & Connectivity	Value For Money
Recommended			

0	NaN	NaN	3.0
yes			
1	NaN	NaN	2.0
no			
2	NaN	NaN	2.0
no			

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23171 entries, 0 to 23170
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	23171 non-null	int64
1	Airline Name	23171 non-null	object
2	Overall_Rating	23171 non-null	object
3	Review_Title	23171 non-null	object
4	Review_Date	23171 non-null	object
5	Verified	23171 non-null	bool
6	Review	23171 non-null	object
7	Aircraft	7129 non-null	object
8	Type Of Traveller	19433 non-null	object
9	Seat Type	22075 non-null	object
10	Route	19343 non-null	object
11	Date Flown	19417 non-null	object
12	Seat Comfort	19016 non-null	float64
13	Cabin Staff Service	18911 non-null	float64
14	Food & Beverages	14500 non-null	float64
15	Ground Service	18378 non-null	float64
16	Inflight Entertainment	10829 non-null	float64
17	Wifi & Connectivity	5920 non-null	float64
18	Value For Money	22105 non-null	float64
19	Recommended	23171 non-null	object

```
dtypes: bool(1), float64(7), int64(1), object(11)
```

```
memory usage: 3.4+ MB
```

```
df.drop("Unnamed: 0",axis=1,inplace=True)
```

```
df.drop("Date Flown",axis=1,inplace=True)
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23171 entries, 0 to 23170
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Airline Name                          23171 non-null  object
1   Overall_Rating                        23171 non-null  object
2   Review_Title                          23171 non-null  object
3   Review_Date                           23171 non-null  object
4   Verified                              23171 non-null  bool
5   Review                                23171 non-null  object
6   Aircraft                              7129 non-null   object
7   Type Of Traveller                     19433 non-null  object
8   Seat Type                             22075 non-null  object
9   Route                                 19343 non-null  object
10  Seat Comfort                           19016 non-null  float64
11  Cabin Staff Service                   18911 non-null  float64
12  Food & Beverages                      14500 non-null  float64
13  Ground Service                       18378 non-null  float64
14  Inflight Entertainment                10829 non-null  float64
15  Wifi & Connectivity                   5920 non-null   float64
16  Value For Money                       22105 non-null  float64
17  Recommended                           23171 non-null  object
dtypes: bool(1), float64(7), object(10)
memory usage: 3.0+ MB

# Checking duplicate rows:
df.duplicated().value_counts()

False    23051
True       120
Name: count, dtype: int64

df = df.drop_duplicates().reset_index(drop=True)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23051 entries, 0 to 23050
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Airline Name                          23051 non-null  object
1   Overall_Rating                        23051 non-null  object
2   Review_Title                          23051 non-null  object
3   Review_Date                           23051 non-null  object
4   Verified                              23051 non-null  bool
5   Review                                23051 non-null  object
6   Aircraft                              7127 non-null   object
7   Type Of Traveller                     19424 non-null  object

```

8	Seat Type	22062	non-null	object
9	Route	19335	non-null	object
10	Seat Comfort	19007	non-null	float64
11	Cabin Staff Service	18902	non-null	float64
12	Food & Beverages	14496	non-null	float64
13	Ground Service	18369	non-null	float64
14	Inflight Entertainment	10827	non-null	float64
15	Wifi & Connectivity	5918	non-null	float64
16	Value For Money	22092	non-null	float64
17	Recommended	23051	non-null	object

dtypes: bool(1), float64(7), object(10)

memory usage: 3.0+ MB

df.duplicated().value\_counts()

False      23051

Name: count, dtype: int64

*# For categorical columns*

df['Aircraft']=df['Aircraft'].fillna(df['Aircraft'].mode().iloc[0])

df['Type Of Traveller']=df['Type Of Traveller'].fillna(df['Type Of Traveller'].mode().iloc[0])

df['Seat Type']=df['Seat Type'].fillna(df['Seat Type'].mode().iloc[0])

df['Route']=df['Route'].fillna(df['Route'].mode().iloc[0])

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 23051 entries, 0 to 23050

Data columns (total 18 columns):


#	Column	Non-Null Count	Dtype
0	Airline Name	23051 non-null	object
1	Overall_Rating	23051 non-null	object
2	Review_Title	23051 non-null	object
3	Review Date	23051 non-null	object
4	Verified	23051 non-null	bool
5	Review	23051 non-null	object
6	Aircraft	23051 non-null	object
7	Type Of Traveller	23051 non-null	object
8	Seat Type	23051 non-null	object
9	Route	23051 non-null	object
10	Seat Comfort	19007 non-null	float64
11	Cabin Staff Service	18902 non-null	float64
12	Food & Beverages	14496 non-null	float64
13	Ground Service	18369 non-null	float64
14	Inflight Entertainment	10827 non-null	float64
15	Wifi & Connectivity	5918 non-null	float64
16	Value For Money	22092 non-null	float64
17	Recommended	23051 non-null	object

```
dtypes: bool(1), float64(7), object(10)
memory usage: 3.0+ MB
```

```
# For numerical columns
l=['Seat Comfort','Cabin Staff Service','Food & Beverages','Ground
Service','Inflight Entertainment','Wifi & Connectivity','Value For
Money']
for i in l:
    df[i]=df[i].fillna(df[i].median())
```

## EDA (Data Visualization):


```
# Verified Reviews
plt.figure(figsize=(4, 4))
sns.countplot(x='Verified', data=df)
plt.title('Number of Verified Reviews')
plt.xlabel('Verified')
plt.ylabel('Count')
plt.show()
```




```
# Type of Traveller Distribution
traveler_distribution = df['Type Of Traveller'].value_counts()

# Distribution of types of Travellers
plt.figure(figsize=(6, 4))
traveler_distribution.plot(kind='pie', autopct='%1.1f%%')
```


```
plt.title('Distribution of Types of Travellers')
plt.ylabel('')
plt.show()
```



```
# Percetage of recommened reviews
recommended_percentage =
df['Recommended'].value_counts(normalize=True) * 100
plt.figure(figsize=(4, 4))
recommended_percentage.plot(kind='bar', rot=0)
plt.title('Percentage of Recommended Reviews')
plt.xlabel('Recommended')
plt.ylabel('Percentage')
plt.show()
```



```
# Histogram of Overall Ratings:
plt.figure(figsize=(6, 4))
sns.histplot(df['Overall_Rating'], bins=10, kde=True)
plt.title('Distribution of Overall Ratings')
plt.xlabel('Overall Rating')
plt.ylabel('Frequency')
plt.show()
```



```
# Bar Plot for Top Routes:
top_routes = df['Route'].value_counts().head(10)
plt.figure(figsize=(5, 4))
colors = sns.color_palette("Set1", len(top_routes))
sns.barplot(x=top_routes.index, y=top_routes.values, palette=colors)
plt.title('Top Routes by Review Count')
plt.xlabel('Route')
plt.ylabel('Review Count')
plt.xticks(rotation=90)
```

C:\Users\Vidul\AppData\Local\Temp\ipykernel\_34468\2586659521.py:5:  
FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_routes.index, y=top_routes.values, palette=colors)

([0, 1, 2, 3, 4, 5, 6, 7, 8, 9],
 [Text(0, 0, 'Melbourne to Sydney'),
  Text(1, 0, 'Sydney to Melbourne'),
  Text(2, 0, 'Cape Town to Johannesburg'),
  Text(3, 0, 'Cusco to Lima'),
  Text(4, 0, 'Bangkok to Phuket'),
  Text(5, 0, 'Johannesburg to Cape Town'),
  Text(6, 0, 'Kuala Lumpur to Singapore'),
```




```
Text(7, 0, 'Bangkok to Chiang Mai'),
Text(8, 0, 'Johannesburg to Durban'),
Text(9, 0, 'Toronto to Calgary')])
```



```
# Pairwise Scatterplot for Correlation:
sns.pairplot(df, vars=['Seat Comfort', 'Cabin Staff Service', 'Value
For Money', 'Overall_Rating'])
```

```
<seaborn.axisgrid.PairGrid at 0x201094c2050>
```



```
## Removing unnecessary columns:
df.drop("Route",axis=1,inplace=True)
df.drop("Aircraft",axis=1,inplace=True)
df.drop("Review Date",axis=1,inplace=True)
df.drop("Airline Name",axis=1,inplace=True)
```

## Encoding

```
df['Verified']=df['Verified'].replace({'True':1,'False':0})
df['Type Of Traveller']=df['Type Of Traveller'].replace({'Solo
Leisure':1,'Couple Leisure':2,'Family Leisure':3,'Business':4})
df['Seat Type']=df['Seat Type'].replace({'Economy Class':1,'Business
```

```

Class':2,'Premium Economy':3,'First Class':4})
df['Recommended']=df['Recommended'].replace({'yes':1,'no':0})

df['Overall_Rating']=df['Overall_Rating'].replace('n',10)
df['Overall_Rating']=df['Overall_Rating'].astype(int)

l=["Overall_Rating","Review_Title","Cabin Staff Service","Food &
Beverages","Recommended"]
df=df[l]

df.head(3)

```

	Overall_Rating	Review_Title	Cabin Staff Service
0	9	"pretty decent airline"	5.0
1	1	"Not a good airline"	2.0
2	1	"flight was fortunately short"	1.0

	Food & Beverages	Recommended
0	4.0	1
1	1.0	0
2	1.0	0

## Preprocessing using spacy:

```

nlp=spacy.load("en_core_web_sm")

C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\spacy\util.py:910: UserWarning: [W095] Model 'en_core_web_sm'
(3.6.0) was trained with spaCy v3.6.0 and may not be 100% compatible
with the current version (3.7.2). If you see errors or degraded
performance, download a newer compatible model or retrain your custom
model with the current spaCy version. For more details and available
updates, run: python -m spacy validate
  warnings.warn(warn_msg)

# Important words (so that they are not lost)
exceptions = ["not", "never", "bad", "nice", "good", "great", "poor",
"excellent","no"]
def preprocess (text):
    text = text.lower() # to convert to lowercase (so that we can add
exceptions properly)
    doc=nlp(text)
    filtered_tokens=[]
    for token in doc:
        if token.text in exceptions:
            filtered_tokens.append(token.text)

```

```

        if token.is_stop or token.is_punct:
            continue
        filtered_tokens.append(token.lemma_)
    return " ".join(filtered_tokens)

```

```
df['Review']=df['Review_Title'].apply(lambda x: preprocess(x))
```

C:\Users\Vidul\AppData\Local\Temp\ipykernel\_34468\4282974474.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['Review']=df['Review_Title'].apply(lambda x: preprocess(x))
```

```
df.head(5)
```

	Overall_Rating	Review_Title	Cabin Staff Service \
0	9	"pretty decent airline"	
1	1	"Not a good airline"	
2	1	"flight was fortunately short"	
3	1	"I will never fly again with Adria"	
4	1	"it ruined our last days of holidays"	

	Food & Beverages	Recommended	Review
0	4.0	1	pretty decent airline
1	1.0	0	not good good airline
2	1.0	0	flight fortunately short
3	2.0	0	never fly adria
4	1.0	0	ruin day holiday

```
df.columns
```

```
Index(['Overall_Rating', 'Review_Title', 'Cabin Staff Service',
      'Food & Beverages', 'Recommended', 'Review'],
      dtype='object')
```

```
df=df.rename(columns={'Cabin Staff Service': 'Cabin_Staff_Service',
                     'Food & Beverages': 'Food_Beverages'})
```

```
v1=TfidfVectorizer()
v1.fit(df["Review"])
rev_tfidf = v1.transform(df["Review"])

rev_tfidf

<23051x3892 sparse matrix of type '<class 'numpy.float64'>'
  with 68754 stored elements in Compressed Sparse Row format>

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23051 entries, 0 to 23050
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Overall_Rating        23051 non-null  int32
 1   Review_Title          23051 non-null  object
 2   Cabin_Staff_Service  23051 non-null  float64
 3   Food_Beverages       23051 non-null  float64
 4   Recommended           23051 non-null  int64
 5   Review                23051 non-null  object
dtypes: float64(2), int32(1), int64(1), object(2)
memory usage: 990.6+ KB
```

[illegible]



```

4          1.0
...      ...
23046     2.0
23047     2.0
23048     2.0
23049     3.0
23050     2.0

[23051 rows x 3895 columns]

x.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23051 entries, 0 to 23050
Columns: 3895 entries, 0 to Food_Beverages
dtypes: float64(3894), int32(1)
memory usage: 684.9 MB

y.info()

<class 'pandas.core.series.Series'>
RangeIndex: 23051 entries, 0 to 23050
Series name: Recommended
Non-Null Count  Dtype
-----
23051 non-null  int64
dtypes: int64(1)
memory usage: 180.2 KB

x.columns = x.columns.astype(str)

# scaler = MinMaxScaler()
# scaler.fit(x)
# x = scaler.transform(x)
# x = pd.DataFrame(x)
# x.info()

```

## TrainTestSplit

```

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)

print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

(18440, 3895)
(4611, 3895)
(18440,)
(4611,)

```

```

y_train
14149    0
3728     1
16249    1
450      0
17987    1
..
11964    0
21575    0
5390     0
860      0
15795    0
Name: Recommended, Length: 18440, dtype: int64

```

## Model Training

### 1) Neural Networks

```

model1 = keras.Sequential([
    keras.layers.Input(shape=(x_train.shape[1],)), # shape of the
input data
    keras.layers.Dense(64, activation='relu'), # hidden layer
    keras.layers.Dense(1, activation='sigmoid')
])

model1.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)

# Train the model
model1.fit(x_train, y_train, epochs=50)

Epoch 1/50
577/577 [=====] - 7s 7ms/step - loss: 0.3107
- accuracy: 0.8885
Epoch 2/50
577/577 [=====] - 4s 7ms/step - loss: 0.1759
- accuracy: 0.9310
Epoch 3/50
577/577 [=====] - 4s 7ms/step - loss: 0.1486
- accuracy: 0.9427
Epoch 4/50
577/577 [=====] - 5s 8ms/step - loss: 0.1335
- accuracy: 0.9494
Epoch 5/50
577/577 [=====] - 5s 8ms/step - loss: 0.1220

```



```
- accuracy: 0.9530
Epoch 6/50
577/577 [=====] - 4s 7ms/step - loss: 0.1120
- accuracy: 0.9584
Epoch 7/50
577/577 [=====] - 4s 8ms/step - loss: 0.1053
- accuracy: 0.9617
Epoch 8/50
577/577 [=====] - 4s 7ms/step - loss: 0.0997
- accuracy: 0.9653
Epoch 9/50
577/577 [=====] - 4s 7ms/step - loss: 0.0943
- accuracy: 0.9670
Epoch 10/50
577/577 [=====] - 4s 7ms/step - loss: 0.0901
- accuracy: 0.9685
Epoch 11/50
577/577 [=====] - 4s 7ms/step - loss: 0.0864
- accuracy: 0.9695
Epoch 12/50
577/577 [=====] - 4s 7ms/step - loss: 0.0829
- accuracy: 0.9716
Epoch 13/50
577/577 [=====] - 4s 7ms/step - loss: 0.0803
- accuracy: 0.9729
Epoch 14/50
577/577 [=====] - 4s 7ms/step - loss: 0.0790
- accuracy: 0.9734
Epoch 15/50
577/577 [=====] - 4s 6ms/step - loss: 0.0757
- accuracy: 0.9739
Epoch 16/50
577/577 [=====] - 5s 8ms/step - loss: 0.0740
- accuracy: 0.9751
Epoch 17/50
577/577 [=====] - 5s 8ms/step - loss: 0.0715
- accuracy: 0.9755
Epoch 18/50
577/577 [=====] - 4s 7ms/step - loss: 0.0704
- accuracy: 0.9765
Epoch 19/50
577/577 [=====] - 5s 8ms/step - loss: 0.0688
- accuracy: 0.9768
Epoch 20/50
577/577 [=====] - 4s 8ms/step - loss: 0.0679
- accuracy: 0.9767
Epoch 21/50
577/577 [=====] - 5s 8ms/step - loss: 0.0672
- accuracy: 0.9758
```

Epoch 22/50  
577/577 [=====] - 5s 8ms/step - loss: 0.0650  
- accuracy: 0.9783  
Epoch 23/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0637  
- accuracy: 0.9786  
Epoch 24/50  
577/577 [=====] - 5s 8ms/step - loss: 0.0623  
- accuracy: 0.9798  
Epoch 25/50  
577/577 [=====] - 5s 8ms/step - loss: 0.0616  
- accuracy: 0.9793  
Epoch 26/50  
577/577 [=====] - 4s 7ms/step - loss: 0.0596  
- accuracy: 0.9805  
Epoch 27/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0598  
- accuracy: 0.9800  
Epoch 28/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0593  
- accuracy: 0.9804  
Epoch 29/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0562  
- accuracy: 0.9818  
Epoch 30/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0553  
- accuracy: 0.9818  
Epoch 31/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0558  
- accuracy: 0.9815  
Epoch 32/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0558  
- accuracy: 0.9813  
Epoch 33/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0536  
- accuracy: 0.9817  
Epoch 34/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0534  
- accuracy: 0.9822  
Epoch 35/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0520  
- accuracy: 0.9826  
Epoch 36/50  
577/577 [=====] - 4s 7ms/step - loss: 0.0510  
- accuracy: 0.9825  
Epoch 37/50  
577/577 [=====] - 4s 8ms/step - loss: 0.0513  
- accuracy: 0.9827  
Epoch 38/50

```
577/577 [=====] - 4s 8ms/step - loss: 0.0511
- accuracy: 0.9828
Epoch 39/50
577/577 [=====] - 4s 8ms/step - loss: 0.0500
- accuracy: 0.9832
Epoch 40/50
577/577 [=====] - 5s 8ms/step - loss: 0.0493
- accuracy: 0.9845
Epoch 41/50
577/577 [=====] - 4s 8ms/step - loss: 0.0493
- accuracy: 0.9826
Epoch 42/50
577/577 [=====] - 4s 8ms/step - loss: 0.0481
- accuracy: 0.9833
Epoch 43/50
577/577 [=====] - 4s 8ms/step - loss: 0.0470
- accuracy: 0.9847
Epoch 44/50
577/577 [=====] - 4s 8ms/step - loss: 0.0468
- accuracy: 0.9832
Epoch 45/50
577/577 [=====] - 4s 8ms/step - loss: 0.0466
- accuracy: 0.9844
Epoch 46/50
577/577 [=====] - 4s 8ms/step - loss: 0.0459
- accuracy: 0.9846
Epoch 47/50
577/577 [=====] - 4s 8ms/step - loss: 0.0460
- accuracy: 0.9844
Epoch 48/50
577/577 [=====] - 4s 8ms/step - loss: 0.0455
- accuracy: 0.9844
Epoch 49/50
577/577 [=====] - 5s 8ms/step - loss: 0.0450
- accuracy: 0.9850
Epoch 50/50
577/577 [=====] - 4s 8ms/step - loss: 0.0434
- accuracy: 0.9858
```

```
<keras.src.callbacks.History at 0x2010b3a69d0>
```

```
# Evaluate the model
```

```
loss, accuracy = model1.evaluate(x_train, y_train)
```

```
print(f"Test accuracy: {accuracy}")
```

```
577/577 [=====] - 5s 4ms/step - loss: 0.0396
- accuracy: 0.9867
Test accuracy: 0.9867136478424072
```

## 2) Random Forest Classifier

```
model2=RandomForestClassifier(n_estimators=100, random_state=42)
model2
```

```
RandomForestClassifier(random_state=42)
```

```
model2.fit(x_train, y_train)
y_pred2 = model2.predict(x_test)
print(classification_report(y_test, y_pred2))
```

	precision	recall	f1-score	support
0	0.95	0.96	0.96	3051
1	0.92	0.90	0.91	1560
accuracy			0.94	4611
macro avg	0.94	0.93	0.93	4611
weighted avg	0.94	0.94	0.94	4611

## 3) KNeighbors Classifier

```
model3=KNeighborsClassifier()
model3
```

```
KNeighborsClassifier()
```

```
model3.fit(x_train, y_train)
y_pred3 = model3.predict(x_test)
print(classification_report(y_test, y_pred3))
```

	precision	recall	f1-score	support
0	0.95	0.95	0.95	3051
1	0.91	0.90	0.91	1560
accuracy			0.94	4611
macro avg	0.93	0.93	0.93	4611
weighted avg	0.94	0.94	0.94	4611

## 4) Decision Tree Classifier

```
model4=DecisionTreeClassifier() # Increase the max_iter value
model4
```

```
DecisionTreeClassifier()
```

```

model4.fit(x_train, y_train)
y_pred4 = model4.predict(x_test)
print(classification_report(y_test, y_pred4))

```

	precision	recall	f1-score	support
0	0.95	0.95	0.95	3051
1	0.90	0.90	0.90	1560
accuracy			0.93	4611
macro avg	0.93	0.93	0.93	4611
weighted avg	0.93	0.93	0.93	4611

## 5) Logistic Regression

```

model5=LogisticRegression(max_iter=1000) # Increase the max_iter value
model5

```

```

LogisticRegression(max_iter=1000)

```

```

model5.fit(x_train, y_train)
y_pred5 = model5.predict(x_test)
print(classification_report(y_test, y_pred5))
print(confusion_matrix(y_test, y_pred5))

```

	precision	recall	f1-score	support
0	0.92	0.93	0.92	3051
1	0.85	0.84	0.85	1560
accuracy			0.90	4611
macro avg	0.89	0.88	0.88	4611
weighted avg	0.90	0.90	0.90	4611

```

[[2827  224]
 [ 254 1306]]

```

```

# Sample input data

```

```

input_data = [['This airline is great', 9, 5, 4], ['Not a good
experience', 2, 3, 2]]

```

```

# Create a DataFrame with the input data

```

```

input_df = pd.DataFrame(input_data, columns=['Review',
'Overall_Rating', 'Cabin Staff Service', 'Food & Beverages'])

```

```

# Preprocess the input data

```

```

# 1. Use the same TfidfVectorizer that you used for training data
rev_tfidf = v1.transform(input_df["Review"])

```

```

# 2. Concatenate the text features and numerical features

```

```

x_input = pd.concat([pd.DataFrame(rev_tfidf.toarray()),
input_df.drop(columns=["Review"])], axis=1)
x_input.columns = x_input.columns.astype(str)

# Predict "Recommended" for the input data
y_pred_input = model1.predict(x_input)

# Selecting 0.5 as a threshold in model1
for i, review in enumerate(input_data):
    if(y_pred_input[i]<=0.5):
        print(f'Review: {review[0]}, Prediction: Not Recommended')
    if(y_pred_input[i]>0.5):
        print(f'Review: {review[0]}, Prediction: Recommended')

1/1 [=====] - 2s 2s/step
Review: This airline is great, Prediction: Recommended
Review: Not a good experience, Prediction: Not Recommended

columns=['Review', 'Overall_Rating', 'Cabin_Staff_Service',
'Food_Beverages']

# Save the trained TF-IDF vectorizer
with open('tfidf_vectorizer.pkl', 'wb') as f:
    pickle.dump(v1, f)

# Save the trained Random Forest classifier
with open('model.pkl', 'wb') as f:
    pickle.dump(model2, f)

# Save the column names of the feature matrix
with open('column_names.pkl', 'wb') as f:
    pickle.dump(columns, f)

```